

# Data Skills Workshop

---

Matthias Haber

3-4 February 2020

# Introduction

---

# Survey results

## pdf

## 2

# About myself

- Head of Data at Looping Studios since 2018
- Postdoc at Hertie 2015-2017 (Governance Report)
- PhD in PolSci (University of Mannheim)
- Research on parties, legislative politics, electoral behavior
- First started programming in 2011

## Contact:

- [haber.matthias@gmail.com](mailto:haber.matthias@gmail.com)

# About yourself

- Who are you?
- Why did you take this class?
- What data/programming skills would make you work life easier?

# Workshop structure

**Day 1** Session 1: A basic introduction to R Session 2: Programming  
Session 3: Tidyverse Session 4: Visualizing trends and relationships

**Day 2** Session 1: Creating our own dataset Session 2: Dashboard  
fundamentals Session 3: Building our own dashboard Session 4:  
Moving deeper and further

# Why data skills?

- Data skills are increasingly important for research and industry projects
- With complex data projects, however, come complex needs for understanding and communicating processes and results

## The 80-20 rule

- Most data are messy
  - You spent most of your time cleaning/preparing data
  - You learn lot about the structure of your data

- Based on the statistical programming language S (1976)
- R was developed by Ross Ihaka and Robert Gentleman (1995)
- R was intentionally developed to be a data analysis language



# Why R

- Open source: makes it highly customizable and easily extensible
- Over 7,500 packages and counting
- Used by many social scientists interested in data analysis
- Powerful tool to generate elegant and effective plots
- Command-line interface and scripts favors reproducibility
- Excellent documentation and online help resources

# We will work in RStudio

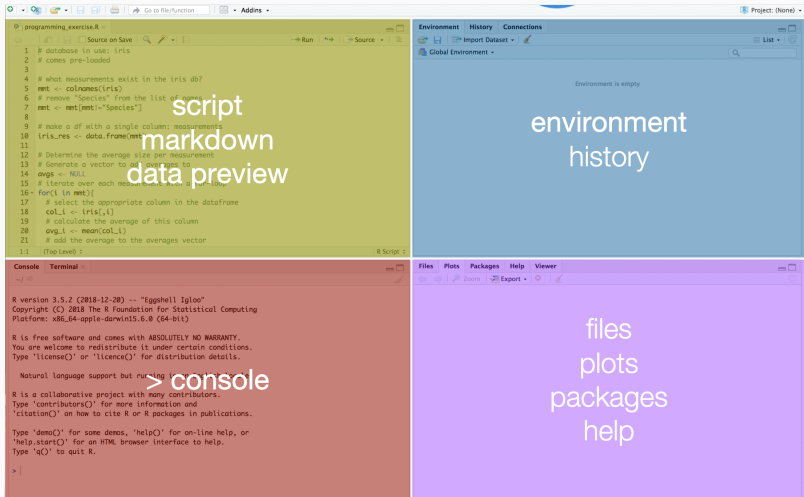
- RStudio is an Integrated Developer Environment (IDE) and serves as:
  - Code editor
    - Code highlighting/completion, indentation, . . .
    - Feed code from editor to R-console
- Project manager
- Workspace viewer
- Data browser
- Enhanced output viewer
- Help browser

- R:
  - R: <http://cran.rstudio.com/>
  - RStudio: <http://www.rstudio.org/download/daily/desktop/>
- R packages
  - tidyverse (R packages designed for data science)

## **Session 1: A basic introduction to R**

---

# The RStudio Interface



- Edit in code editor (.r-file)
- Paste to console
- Save Workspace/Datasets (.Rdata-file)
- Save code routinely (no auto-save!)
- Press TAB to use RStudio's autocompletion feature

- Run code from editor: Select line and `ctrl+Enter`
- Switch between source and console: `ctrl+1`, `ctrl+2`
- Clear console: `ctrl+L`
- 'Arrow up' gives you the last line of code in the console
- Press `Alt+Shift+K` to see all keyboard shortcuts

# Fundamentals of the R language

- Use `#` to comment code (will not be run)
- Case-sensitivity: `data` vs `Data`
- Assigning objects: `<-` and `=`

*# Assign the number 5 to an object called number*

```
number <- 5
```

```
number
```

```
## [1] 5
```

*# Assign the character string Hello World*

```
string <- "Hello World"
```

```
string
```

```
## [1] "Hello World"
```



# Naming

- object names must start with a letter, and can only contain letters, numbers, \_ and ..
- object names should be descriptive
- Each object name must be unique in an environment.
  - Assigning something to an object name that is already in use will overwrite the object's previous contents.

`i_use_snake_case`

`otherPeopleUseCamelCase`

`some.people.use.periods`

`And_aFew.People_RENOUNCEconvention`

- Functions perform operations on the input given and end in ()
- R has a large collection of built-in functions that are called like this:

- For example, `seq()` which makes regular **sequences** of numbers

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

## Operations on scalars

You can use R as a calculator:

```
2 + 3
```

```
2 - 3
```

```
2 * 3
```

```
2 / 3
```

Functions on scalars:

```
a <- 5
```

```
factorial(a)
```

```
## [1] 120
```

## Exercise 1

1. Do the following calculation in R:  $\frac{1+5}{9}$
2. Assign the results to a variable
3. Bonus: Round off the results to the 1 decimal

## Special values in R

- NA: not available, missing
- NULL: does not exist, is undefined
- TRUE, T: logical true
- FALSE, F: logical false

## Finding special values

Function	Meaning
<code>is.na</code>	Is the value NA
<code>is.null</code>	Is the value NULL
<code>isTRUE</code>	Is the value TRUE
<code>!isTRUE</code>	Is the value FALSE

```
absent <- NA  
is.na(absent)  
  
## [1] TRUE
```

Operator	Meaning
<	less than
>	greater than
==	equal to
<=	less than or equal to
>=	greater than or equal to
!=	not equal to
a   b	a or b
a & b	a and b



# R is object-oriented

Objects are R's nouns and include (not exhaustive):

- character strings
- numbers
- vectors of numbers or character strings
- matrices
- data frames
- lists

# Vectors

A vector is a container of objects put together in an order.

```
# Define a vector
```

```
a <- c(1,4,5)
```

```
b <- c(3,6,7)
```

```
# Join multiple vectors
```

```
ab <- c(a,b)
```

```
# Find vector length (number of its elements)
```

```
length(a)
```

# Operations on vectors

Operation	Meaning
<code>sort(x)</code>	sort a vector
<code>sum(x)</code>	sum of vector elements
<code>mean(x)</code>	arithmetic mean
<code>median(x)</code>	median value
<code>var(x)</code>	variance
<code>sd(x)</code>	standard deviation
<code>factorial(x)</code>	factorial of a number

## Exercise 2

1. Create a character vector with the names of the three people sitting closest to you. Save the vector as `name`
2. Create a numeric vector with their respective ages and save it as `age`
3. Use a function in R to calculate their average age?

A Matrix is a square 2 dimensional container, i.e. vectors combined by row or column

- Must specify number of rows and columns

`matrix(x,nrow,ncol,byrow)`

- `x`: vector of length `nrow*ncol`
- `nrow`: number of rows
- `ncol`: number of columns
- `byrow`: TRUE or FALSE, specifies direction of input

## Exercise 3

Assign a  $6 \times 10$  matrix with the sequence  $1, 2, 3, \dots, 60$  as the data.  
Save the matrix as `m`

# Data frames

Data frames are a two-dimensional container of vectors with the same length. Each column (vector) can be of a different class and can be referenced or created with \$. You can use functions like `nrow()`, `ncol()`, `dim()`, `colnames()`, or `rownames()` on your `df`.

*# Combine two vectors into a data frame*

```
number <- c(1, 2, 3, 4)
```

```
name <- c('John', 'Paul', 'George', 'Ringo')
```

```
df <- data.frame(number, name, stringsAsFactors = FALSE)
```

```
df
```

```
##   number   name
```

```
## 1      1   John
```

```
## 2      2   Paul
```

```
## 3      3 George
```

```
## 4      4  Ringo
```

## Exercise 4

1. Create a vector called `country` containing the names of the countries from the three people whose names you used earlier.
2. Create a data frame combining `name`, `age` and `country` and save it as `my_first_df`



# Lists

A list is an object containing other objects that can have different lengths and classes.

```
# Create a list with three objects of different lengths
list1 <- list(beatles = c('John', 'Paul', 'George', 'Ringo'
                        alive = c('Paul', 'Ringo'), albums = 1:13)

list1

## $beatles
## [1] "John"    "Paul"    "George"  "Ringo"
##
## $alive
## [1] "Paul"    "Ringo"
##
## $albums
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13
```

## Exercise 5

1. Add one more person's name to `name` vector
2. Try to create a data frame called `my_second_df` and store the new `name` vector, `age` and `country` in it. See what happens and why.
3. Create a list instead of a data frame with the three objects and name it `my_first_list`

# Indexing vectors

## By Position

<code>x[4]</code>	The fourth element.
<code>x[-4]</code>	All but the fourth.
<code>x[2:4]</code>	Elements two to four.
<code>x[-(2:4)]</code>	All elements except two to four.
<code>x[c(1, 5)]</code>	Elements one and five.

## By Value

<code>x[x == 10]</code>	Elements which are equal to 10.
<code>x[x &lt; 0]</code>	All elements less than zero.
<code>x[x %in% c(1, 2, 5)]</code>	Elements in the set 1, 2, 5.

## Named Vectors

<code>x['apple']</code>	Element with name 'apple'.
-------------------------	----------------------------

## Exercise 6

1. Return the first number in your vector `age`
2. Return the 2nd and 3rd element in your vector `name`
3. Return only ages under 30 from your vector `age`

# Referencing matrices

- Like vectors, you can reference matrices by elements
- Can also reference rows/columns, these are vectors



`m[2, ]` - Select a row



`m[ , 1]` - Select a column



`m[2, 3]` - Select an element

## Exercise 7

Extract the 9th column of the matrix from the previous problem.  
How can you find the 4th element in the 9th column?

## Indexing data frames

`df[ , 2]`




`df[2, ]`



`df[2, 2]`



`df$x`



## Exercise 8

1. From your data frame `my_first_df`, return the entries for everyone living in a country of your choice.



## Indexing lists

```
list1[1] # 1st element of the list
```

```
## $beatles
```

```
## [1] "John"    "Paul"    "George"  "Ringo"
```

```
list1[[1]] # 1st content of the first element
```

```
## [1] "John"    "Paul"    "George"  "Ringo"
```

```
list1[[1]][[2]] # 2nd value in the 1st content
```

```
## [1] "Paul"
```

- `[]` for indexing vectors, lists, data frames. . .
- `()` for passing arguments to functions
- `{ }` for defining content of loops, functions, etc.

# Recap types and structures

- Data types encountered so far:
  - logical
  - numeric
  - character
- Data structures
  - vector (1 dimension)
  - matrix' (2 dimensions)
  - data frame (2 dimensions)
  - list ( $n$  dimensions)
- Absent data
  - NA (not available)
  - NULL (non-existent)

# Recap functions

- Functions encountered so far
  - `c()`
  - `data.frame()`
  - `mean()`
  - ...
- What if you don't know what a functions does?
  - `?mean()` to get help for a function
  - `help.search('weighted mean')` to get help for a concept

```
mean {base}
```

## Arithmetic Mean

### Description

Generic function for the (trimmed) arithmetic mean.

### Usage

```
mean(x, ...)
```

```
## Default S3 method:
```

```
mean(x, trim = 0, na.rm = FALSE, ...)
```

### Arguments

- x** An R object. Currently there are methods for numeric/logical vectors and [date](#), [date-time](#) and [time interval](#) objects. Complex vectors are allowed for `trim = 0`, only.
- trim** the fraction (0 to 0.5) of observations to be trimmed from each end of `x` before the mean is computed. Values of `trim` outside that range are taken as the nearest endpoint.
- na.rm** a logical value indicating whether NA values should be stripped before the computation proceeds.
- ...** further arguments passed to or from other methods.

## Value

If `trim` is zero (the default), the arithmetic mean of the values in `x` is computed, as a numeric or complex vector of length one. If `x` is not logical (coerced to numeric), numeric (including integer) or complex, `NA_real_` is returned, with a warning.

If `trim` is non-zero, a symmetrically trimmed mean is computed with a fraction of `trim` observations deleted from each end before the mean is computed.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

[weighted.mean](#), [mean.POSIXct](#), [colMeans](#) for row and column means.

## Examples

```
x <- c(0:10, 50)
xm <- mean(x)
c(xm, mean(x, trim = 0.10))
```

## R's build-in data sets

There are a number of example data sets available within R.

```
# List internal data sets:
```

```
data()
```

```
# Load swiss data set:
```

```
data(swiss)
```

```
# Find data description:
```

```
?swiss
```

## **Session 2: Messing with data**

---



- Complete the (free) course on introduction to R:  
<https://www.datacamp.com/courses/free-introduction-to-r>

**That's it for today. Questions?**